

Assignment 6 – Factor Analysis

Project Name: Prediction of sales prices of houses

GitHub link: <https://github.com/tnutalapati/prediction-of-sales-prices-of-houses>

Team:

Tejaswini Nutalapati

Aditi Bhargava

Loading the libraries

```
library(ggplot2)
library(caret)
library(scales)
library(dummies)
library(fmsb)
library(randomForest)
library(corrplot)
library(ggrepel)
library(rlang)
library(gridExtra)
library(knitr)
library(Amelia)
library(mice)
library(dplyr)
library(fastDummies)
library(lattice)
library(pcaPP)
```

head(train,2)

```
head(train,2)
# A tibble: 2 x 82
  Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape LandContour Utilities LotConfig LandSlope Neighborhood
<dbl> <dbl> <chr> <chr> <dbl> <chr> <fct> <chr> <chr> <chr> <chr> <chr>
1 1 60 RL 65 8450 Pave NA Reg Lvl AllPub Inside Gtl CollgCr
2 2 20 RL 80 9600 Pave NA Reg Lvl AllPub FR2 Gtl Veenker
... with 69 more variables: Condition1 <chr>, Condition2 <chr>, BldgType <chr>, HouseStyle <chr>, OverallQual <dbl>,
OverallCond <dbl>, YearBuilt <dbl>, YearRemodAdd <dbl>, RoofStyle <chr>, RoofMatl <chr>, Exterior1st <chr>, Exterior2nd <chr>,
MasVnrType <fct>, MasVnrArea <chr>, ExterQual <chr>, ExterCond <chr>, Foundation <chr>, BsmtQual <fct>, BsmtCond <fct>,
BsmtExposure <fct>, BsmtFinType1 <fct>, BsmtFinSF1 <chr>, BsmtFinType2 <fct>, BsmtFinSF2 <chr>, BsmtUnfSF <chr>,
TotalBsmtSF <chr>, Heating <chr>, HeatingQC <chr>, CentralAir <chr>, Electrical <chr>, '1stFlrSF' <dbl>, '2ndFlrSF' <dbl>,
LowQualFinSF <dbl>, GrLivArea <dbl>, BsmtFullBath <chr>, BsmtHalfBath <chr>, FullBath <dbl>, HalfBath <dbl>, BedroomAbvGr <dbl>,
KitchenAbvGr <dbl>, KitchenQual <chr>, TotRmsAbvGrd <dbl>, Functional <chr>, Fireplaces <dbl>, FireplaceQu <fct>,
GarageType <fct>, GarageYrBlt <chr>, GarageFinish <fct>, GarageCars <dbl>, GarageArea <dbl>, GarageQual <fct>, GarageCond <fct>,
PavedDrive <chr>, WoodDeckSF <dbl>, OpenPorchSF <dbl>, EnclosedPorch <dbl>, '3SsnPorch' <dbl>, ScreenPorch <dbl>,
PoolArea <dbl>, PoolQC <fct>, Fence <fct>, MiscFeature <fct>, MiscVal <dbl>, MoSold <dbl>, YrSold <dbl>, SaleType <chr>,
SaleCondition <chr>, SalePrice <dbl>, Remodel_flag <fct>
```

```
head(test,2)
```

```
head(test,2)
```

```
A tibble: 2 x 82
```

```
  Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape LandContour Utilities LotConfig LandSlope Neighborhood
<dbl> <dbl> <chr> <chr> <dbl> <chr> <fct> <chr> <chr> <chr> <chr> <chr>
1461 20 RH 80 11622 Pave NA Reg Lvl AllPub Inside Gtl NAmes
1462 20 RL 81 14267 Pave NA IR1 Lvl AllPub Corner Gtl NAmes
... with 69 more variables: Condition1 <chr>, Condition2 <chr>, BldgType <chr>, HouseStyle <chr>, OverallQual <dbl>,
OverallCond <dbl>, YearBuilt <dbl>, YearRemodAdd <dbl>, RoofStyle <chr>, RoofMatl <chr>, Exterior1st <chr>, Exterior2nd <chr>,
MasVnrType <fct>, MasVnrArea <chr>, ExterQual <chr>, ExterCond <chr>, Foundation <chr>, BsmtQual <fct>, BsmtCond <fct>,
BsmtExposure <fct>, BsmtFinType1 <fct>, BsmtFinSF1 <chr>, BsmtFinType2 <fct>, BsmtFinSF2 <chr>, BsmtUnfSF <chr>,
TotalBsmtSF <chr>, Heating <chr>, HeatingQC <chr>, CentralAir <chr>, Electrical <chr>, '1stFlrSF' <dbl>, '2ndFlrSF' <dbl>,
LowQualFinSF <dbl>, GrLivArea <dbl>, BsmtFullBath <chr>, BsmtHalfBath <chr>, FullBath <dbl>, HalfBath <dbl>, BedroomAbvGr <dbl>,
KitchenAbvGr <dbl>, KitchenQual <chr>, TotRmsAbvGrd <dbl>, Functional <chr>, Fireplaces <dbl>, FireplaceQu <fct>,
GarageType <fct>, GarageYrBlt <chr>, GarageFinish <fct>, GarageCars <dbl>, GarageArea <dbl>, GarageQual <fct>, GarageCond <fct>,
PavedDrive <chr>, WoodDeckSF <dbl>, OpenPorchSF <dbl>, EnclosedPorch <dbl>, '3SsnPorch' <dbl>, ScreenPorch <dbl>,
PoolArea <dbl>, PoolQC <fct>, Fence <fct>, MiscFeature <fct>, MiscVal <dbl>, MoSold <dbl>, YrSold <dbl>, SaleType <chr>,
SaleCondition <chr>, SalePrice <dbl>, Remodel_Flag <fct>
```

```
salePrice <- train$SalePrice
train2<-train[-81]
data <- rbind(train2, test)
id<-data[1]
```

```
sum(is.na(data))
```

There is a total of 13965 missing values. Now, checking which variables have missing values

```
i <- 1
na <- 1
for (i in 1:length(data))
{
  na[i] <- ifelse(sum(is.na(data[i]))>0, sum(is.na(data[i])), 0)
  if (na[i]>0)
  {
    cat(names(data[i]), '=', na[i], ' ')
  }
}
```


```
> MSZoning = 4
LotFrontage = 486
Alley = 2721
Utilities = 2
Exterior1st = 1
Exterior2nd = 1
MasVnrType = 24
MasVnrArea = 23
BsmtQual = 81
BsmtCond = 82
BsmtExposure = 82
BsmtFinType1 = 79
BsmtFinSF1 = 1
BsmtFinType2 = 80
BsmtFinSF2 = 1
BsmtUnfSF = 1
TotalBsmtSF = 1
Electrical = 1
BsmtFullBath = 2
BsmtHalfBath = 2
KitchenQual = 1
Functional = 2
FireplaceQu = 1420
GarageType = 157
GarageYrBlt = 159
GarageFinish = 159
GarageCars = 1
GarageArea = 1
GarageQual = 159
GarageCond = 159
PoolQC = 2909
Fence = 2348
MiscFeature = 2814
SaleType = 1 |
```

Variables with about or more than 50% missing data in train and test set is PoolQC, MiscFeature, Alley, Fence and FireplaceQu

Mapping the missingness Map:

```
missmap(data)
```

Miasingnes« uap



```
z<- c(WI ch(coin ees{dotog==" PoolgC" }, sht ch(co1nomes{doteg=="Ht scFeature" }
, nhichCco1nomesCdoto}=="41ley" , wht Ch(ColnorreS{ dotoJ==" Fence"
, whtch(cotnomes£doto>--"Ft reptocegu" ))
```

```
//Removt ng vor l obl es wth mostly nt ssh ng values
doto<-data , c( - 73, - 75, -7, - 74, -38)a
```

```
#Predt ct tve meon matchi ng tregut ott on for numert c , btnory and ordered vort obt es
mt ce_zx>d <- mt ce(doto [ , c( "8smtg uol " , "BsM£ond" , "BsWExposure", "BsmtFt nType1" , " Bsmt FtnType2"
, "E \ ectrt cot " , "GorogeType" , "GorogeFtn\ sh" , "Go regequal " , "GoregeCond"
, "MosVnrType " , "MSZoni ng™" , "UIi T l tt es™" , " E xterl or 1st " , " E xter l or Znd"
, "K tchenquol " , "Funct Sonal " , "SoleType " , "Lotr montage " , "forage YrBt l"
, "NosVnrAreo" , "BsntF tnSF1" , "BsaitF l nSF2" , "BsmJ-Un£ SF" , "Totat BsmrSF"
, "Bsmt Fu\1Both" , "BsnttHol fBoth" , "£orogeCars" , "£orogeAreo“g§ , meLhod- ' pms }
mi ce_conpt ete • - couple tetni ce_niod)
```

Data set preparation

```
#Two sets - numeric and factor
data_factor <- select_if(data, is.factor)
data_numeric<- select_if(data, is.numeric)
data_numeric <- data_numeric[-1]
```

```
#Feature scaling numeric variables
#if PCA is not used there is no need to scale variables in regression tree model)
data_numeric_scale<-data.frame(scale(data_numeric))
```

```
#Creating dummy variables (if PCA is not used there is no need to make dummy variables)
data_factor_dummy <- dummy_cols(data_factor, remove_first_dummy = TRUE)
data_factor_dummy <- select_if(data_factor_dummy, is.numeric)
```

```
data2 <- cbind(id, data_numeric_scale, data_factor_dummy)
```

```
#removing highly correlated variables
datacor<-cor(data2)
f<-findCorrelation(datacor, cutoff = 0.9, verbose = T, names=F)
```

```
data2 <- data2[ ,~f]
```

```
Combination row 106 and column 108 is above the cut-off, value = -0.939
Flagging column 106
Combination row 122 and column 136 is above the cut-off, value = 0.983
Flagging column 122
Combination row 125 and column 139 is above the cut-off, value = 0.97
Flagging column 125
Combination row 129 and column 144 is above the cut-off, value = 0.978
Flagging column 129
Combination row 226 and column 233 is above the cut-off, value = 0.987
Flagging column 226
```

```
#train
train_s <- data2[1:1460,]
train_set<- cbind(train_s, SalePrice)
train_set <- train_set[-1]
#test
test_set <- data2[1461:2919,]
test_set<-test_set[-1]
```

Removing highly correlated variables

```
datacorforest<-cor(data_numeric)
findCorrelation(datacorforest, cutoff = 0.9, verbose = T, names=F)
```

```
> All correlations <= 0.9 |
```

```
train_sett<-data[1:1460,]
train_settt<- cbind(train_sett, SalePrice)
train_settt <- train_settt[-1]
#test no pca
test_settt <- data[1461:2919,]
test_settt<-test_settt[-1]
```

PRINCIPAL COMPONENT ANALYSIS – PCA

Remove dependent variable from train set

```
dim(train_set)
```

1460. 228

Remove dependent variable from train set

```
train_set<-train_set[-228]
```

PCA works only on numeric variables. All values are numerical.

By default, it centers the variable to have mean equals to zero.

With parameter scale. = T, normalize the variables to have standard deviation equals to 1.

```
pca <- prcomp(train_set, scale. = T)
```

center and scale refer to mean and standard deviation of the variables

```
names(pca)
```

'sdev' 'rotation' 'center' 'scale' 'x'

Each column of rotation matrix contains the principal component loading vector. Principal components and first 3 rows

```
pca$rotation[1:3,1:4]
```

	PC1	PC2	PC3	PC4
MSSubClass	-0.00133405	0.1272953	-0.0404936	0.14172673
LotFrontage	0.06587685	-0.1082165	0.1915191	-0.03903983
LotArea	0.03631721	-0.1118286	0.1874054	-0.05478405


The matrix x has the principal component score vectors in a 1460×227 dimension.

```
dim(pca$x)
```

1460 227

PC1 vs PC2 with loadings

```
autoplot(pca, loadings = TRUE)
```




```
#standard deviation of each principal component
std_dev <- pca$sdev
#variance
variance <- std_dev^2
```

divide the variance by sum of total variance -> to compute the proportion of variance explained by each component

```
variance_prop <- variance/sum(variance)
#first principal component explains 6.98% of the variance, second 3.2%, third 2.5%
variance_prop[1:10]
```


```
0.0688016033899863 0.0316744090851641 0.0252170776030978 0.0208644460252106
0.0173608341666115 0.0154386797614029 0.0147949460535524 0.0139073769471355
0.0129561906693511 0.0129149176043199
```

```
#scree plot - the percentage of variance explained by each principal component
plot(variance_prop, xlab = "Principal Component", ylab = "Proportion of Variance Explained",
     , type = "b", xlim=c(0, 20))
```



Cumulative variance plot

```
# ~ 50 components explains around 60% variance in the data set.
plot(cumsum(variance_prop), xlab = "Principal Component",
     ylab = "Cumulative Proportion of Variance Explained",
     type = "b", xlim=c(0, 50))
```




```

#add a column
test_set$SalePrice <- 1
#new training set with principal components
train_set_pca <- data.frame(SalePrice = train$SalePrice, pca$x)
# first 50 PCAs
train_set_pca <- train_set_pca[,1:51]

```

SalePrice	PC 1	PC2	PC 3	PC4	PC 3	PC6	PC 7	PC8
int	< dbl	<dbl>	< 0D>	<dbl>	< dbl	<dbl>	< dbl	<dbl>
70 8 500	4.<2877471	1.6140i7	1.3619S3	1.5 3<4 467	1.068D6910	1.00S508328	1.2 17208	1.2 78a 7 655
2 18 1500	0.4 9S675d 2	-0.067224	-0.2\$ 3 76 B	-0.38 2 3131	-0.7698 8 64	-1.22 037 33 7	-0.1B 331d	-0.06163059

Factor analysis

```

(eigen_data <- round(euroemp_pca$data^2,2))
names(eigen_data) <- paste("PC",1:9,sep="")
eigen_data

sumlambdas <- sum(eigen_data)
sumlambdas

cumvar_data <- cumsum(propvar)
propvar <- round(eigen_data/sumlambdas,2)
propvar

cumvar_data <- cumsum(propvar)
cumvar_data


matlambdas <- rbind(eigen_data,propvar,cumvar_data)
matlambdas

rownames(matlambdas) <- c("Eigenvalues","Prop. variance","Cum. prop. variance")
rownames(matlambdas)

eigvec.emp <- data_pca$rotation
print(data_pca)

```

Very Simple Structure



```

# Taking the first four PCs to generate linear combinations for all the variables with four factors
pcafactors.emp <- eigvec.emp[,1:4]
pcafactors.emp

# Multiplying each column of the eigenvector's matrix by the square-root of the corresponding eigenvalues
unrot.fact.emp <- sweep(pcafactors.emp,MARGIN=2,euroemp_pca$sdev[1:4],`*`)
unrot.fact.emp

# Computing communalities
communalities.emp <- rowSums(unrot.fact.emp^2)
communalities.emp


# Performing the varimax rotation.
#The default in the varimax function is norm=TRUE thus, Kaiser normalization is carried out
rot.fact.emp <- varimax(unrot.fact.emp)
View(unrot.fact.emp)
rot.fact.emp

# The print method of varimax omits loadings less than abs(0.1).
#In order to display all the loadings, it is necessary to ask explicitly the contents of the object $loadings
fact.load.emp <- rot.fact.emp$loadings[1:9,1:4]
fact.load.emp

# Computing the rotated factor scores for the 30 European Countries.
#Notice that signs are reversed for factors F2 (PC2), F3 (PC3) and F4 (PC4)
scale.data <- scale(data[-1])
scale.data
as.matrix(scale.emp)%*%fact.load.emp%%solve(t(fact.load.emp)%*%fact.load.emp)

```

Non Graphical Solutions to Scree Test



```

fit.pc <- principal(data[-1], nfactors=4, rotate="varimax")
fit.pc
round(fit.pc$values, 3)
fit.pc$loadings

# Loadings with more digits
for (i in c(1,3,2,4)) { print(fit.pc$loadings[[1,i]])}

# Communalities
fit.pc$communality

# Rotated factor scores, Notice the columns ordering: RC1, RC3, RC2 and RC4
fit.pc$scores

# Play with FA utilities
fa.parallel(data[-1])
fa.plot(fit.pc)
fa.diagram(fit.pc)
vss(data[-1])

```